

# Rewriting the Parse API in Go

GopherCon 2015, Denver

8 July 2015

Abhishek Kona

Software Engineer at Parse and Facebook

# What is this talk about?


- Why we rewrote the Parse API in Go
- Tools and libraries we built

# What is Parse?



- Backend as a service
- SDKs for iOS, Android, JS, React, Windows, PHP
- Acquired by Facebook in 2013

# Rome was not built in a day - they did not use Parse



# Parse circa 2013

- ~60K apps
- 10 engineers
- Ruby on Rails app



# Scalability Issues in 2013

- Single popular app could take down Parse
- Fixed-size unicorn pool
- Lengthy deploy times
- Spooky action at a distance

# Solution: Rewrite in Go!

Joel on Software

## Things You Should Never Do, Part I

*by Joel Spolsky*

VS



**CODING HORROR**

programming and human factors

Google™ Custom Search



18 Sep 2006

## When Understanding means Rewriting

# Why Rewrite?

- Huge estimated reliability win
- Wanted easier deploys
- Needed faster tests
- Hard to evolve to existing Ruby codebase



# Why Go?

- Statically typed
- Good concurrency support
- Dynamic number of worker goroutines per HTTP server
- Easy to attract engineers

# Rules of the Rewrite

- Don't break backwards compatibility
- No downtime

# Initial Ports

- Hosting server
- Parse Push Notification Service (PPNS)
- Maintains long-lived push sockets with Android clients
- Concurrent conns per node increased from 250K to 1.5M

# Mongo Proxy

[github.com/facebookgo/dvara](https://github.com/facebookgo/dvara) (<https://github.com/facebookgo/dvara>)

- Mongo used to limit max number of connections to 20000
- We wrote our own proxy for Mongo in Go
- Made easy by Go runtime's use of non-blocking I/O

# Rollout

- Migrate endpoints one by one
- Dified responses between old and new code using a shadow cluster
- Started with low-traffic read only endpoints
- Graduated to write endpoints

# Comment Goldmine

```
// Note: an unset cache version is treated by ruby as "".  
// Because of this, dirtying this isn't as simple as deleting it - we need to  
// actually set a new value.  
  
// This byte sequence is what ruby expects.  
// yes that's a paren after the second 180, per ruby.  
  
// Inserting and having an op is kinda weird: We already know  
// state zero. But ruby supports it, so go does too.  
  
// single geo query, don't do anything. stupid and does not make sense  
// but ruby does it. Changing this will break a lot of client tests.  
// just be nice and fix it here.  
  
// Ruby sets various defaults directly in the structure and expects them to appear in cache.  
// For consistency, we'll do the same thing.
```

# A Young Language

- Some good libraries: mgo, memcache, etc.
- Some missing libraries

# Libraries / Tools



# Dependency Injection

- Helps instantiate implementations for test and production
- Easy to miss passing a dependency to a struct

# Introducing Inject

[github.com/facebookgo/inject](http://github.com/facebookgo/inject) (<http://github.com/facebookgo/inject>)

- Only occurs at process startup for singletons
- Dependencies declared using struct tags
- Fail instead of guessing

# Dependency Injection Example

```
type Handler struct {
    Scribe *scribe.Client `inject:""`
    Log     logger.Logger    `inject:""`
}

// ServeHTTP a sample implementation
func (h *Handler) ServeHTTP(w ResponseWriter, r *Request) {
    params := extractParams(r)
    h.Scribe.Log(params)
    h.Log("everything ok")
    w.Write(res)
}
```

# Main for Inject

```
func main() {  
  
    var g inject.Graph  
    err := g.Provide(  
        &inject.Object{Value: scribe.NewHTTPScrubClient()},  
        &inject.Object{Value: parse.NewLogger()},  
    )  
    if err != nil {  
        fmt.Fprintln(os.Stderr, err)  
        os.Exit(1)  
    }  
  
    if err := g.Populate(); err != nil {  
        fmt.Fprintln(os.Stderr, err)  
        os.Exit(1)  
    }  
    // rest of main  
}
```

# Initializing and Destroying Injected Objects

[github.com/facebookgo/startstop](http://github.com/facebookgo/startstop) (<http://github.com/facebookgo/startstop>)

- Traverses object graph
- At startup: calls `Start` on each injected object in dependency order
- At shutdown: calls `Stop` on each injected object in reverse dependency order
- Fails on cycles

# Start-Stop Example

```
type ScribeClient {
    Thrift    *ThriftPool `inject:""`
}

// ScribeClient start will be called after ThriftPool.Start
func (s *ScribeClient) Start() error {
    fmt.Println("starting scribe client")
    return nil
}

type ThriftPool struct {
    //
}

// ThriftPool start will be called first.
func (t *ThriftPool) Start() error {
    fmt.Println("starting thrift pool")
    return t.tcpDial()
}

func (t *ThriftPool) Stop() error {
    fmt.Println("stopping thrift pool")
    return t.tcpCloseAll()
}
```

# Graceful Restarts

[github.com/facebookgo/grace](https://github.com/facebookgo/grace) (<https://github.com/facebookgo/grace>)

- Restart servers gracefully on deploys
- On USR2, spawns new process and hands off listening socket

# Error Reporting

[github.com/facebookgo/stackerr](https://github.com/facebookgo/stackerr) (<https://github.com/facebookgo/stackerr>)

- Wrap error calls with stackerr
- Aggregate errors based on stack trace in an in-house system called Logview



# Stackerr Example

```
func main() {  
    err := err2()  
    fmt.Println(err)  
}  
  
func err2() error {  
    err := err1()  
    if err != nil {  
        return stackerr.Wrap(err)  
    }  
    return nil  
}  
  
func err1() error {  
    return stackerr.Wrap(errors.New("failure"))  
}
```

# Stackerr Output

```
failure
/private/tmp/stackerr.go:25          err1
/private/tmp/stackerr.go:17          err2
/private/tmp/stackerr.go:12          main
/usr/local/Cellar/go/1.4.2/libexec/src/runtime/proc.go:72    main
/usr/local/Cellar/go/1.4.2/libexec/src/runtime/asm_amd64.s:2233 goexit
(Stack 2)
/private/tmp/stackerr.go:19          err2
/private/tmp/stackerr.go:12          main
/usr/local/Cellar/go/1.4.2/libexec/src/runtime/proc.go:72    main
/usr/local/Cellar/go/1.4.2/libexec/src/runtime/asm_amd64.s:2233 goexit
```

# Muster

[github.com/facebookgo/muster](https://github.com/facebookgo/muster) (<https://github.com/facebookgo/muster>)

- A library to perform operations in a batch
- Two tunables: MaxBatchSize and BatchTimeout

# Generics

[github.com/facebookgo/generics](https://github.com/facebookgo/generics) (<https://github.com/facebookgo/generics>)


# More Libraries

[github.com/facebookgo](https://github.com/facebookgo) (<https://github.com/facebookgo>)

- Many more small libraries
- httpcontrol, ensure, stack

# We Love Go

[github.com/daaku](https://github.com/daaku) ([github.com/daaku](https://github.com/daaku))



# Results


# Results

- ~175k LOC in Go vs ~130k LOC in Ruby
- ~3 minutes to run all the unit tests (down from 25min)
- Apps start in seconds instead of minutes
- Downsized API server pool by 90%
- Rolling restarts dropped from 30 minutes to 3 minutes



# Parse circa 2015

- >500K apps built on Parse
- 2-3x YoY traffic growth
- Primarily a Go stack



# Observations

- Rewrites are hard
- ~4 engineers over 2 years

# Go Side Effects

- Deploying with static binaries is easy
- Developers are responsible for deploys, not ops

# Recap



<http://tiny.cc/parsego>

# Thank you

Abhishek Kona

Software Engineer at Parse and Facebook

[abhishekk@fb.com](mailto:abhishekk@fb.com) (<mailto:abhishekk@fb.com>)

<http://sheki.in> (<http://sheki.in>)

[@sheki](http://twitter.com/sheki) (<http://twitter.com/sheki>)

